


Save 30% or More With Our Ebook Bundles



Machine Learning Mastery

Making Developers Awesome at Machine Learning

How to Develop Multi-Output Regression Models with Python


[Click to Take the FREE Ensemble Learning Crash-Course](#)


Stacking Ensemble Machine Learning With Python



How to Develop a Random Forest Ensemble in Python

by **Jason Brownlee** on April 20, 2020 in **Ensemble Learning**



How to Develop Voting Ensembles With Python

Share

Share

Last Updated on April 27, 2021



One-vs-Rest and One-vs-One for Multi-Class Classification

It is perhaps the most popular and widely used machine learning algorithm given its good or excellent performance across a wide range of classification and regression predictive modeling problems.

Loving the Tutorials?

It is also easy to use given that it has few key hyperparameters and sensible heuristics for configuring these hyperparameters. [The Ensemble Learning With Python EBook](#) is where you'll find the **Really Good** stuff.

In this tutorial, you will learn how to develop a random forest ensemble for classification and regression.

>> SEE WHAT'S INSIDE

After completing this tutorial, you will know:

- Random forest ensemble is an ensemble of decision trees
- How to use the random forest ensemble for classification
- How to explore the effect of random forest model hyperparameters

Kick-start your project with my new book [Ensemble Learning With Python](#) by-step tutorials and the *Python source code* files.

Let's get started.

- **Update Aug/2020:** Added a common question and answer section

Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE


Never miss a tutorial:



Presented by:




How to Develop Multi-Output Regression Models with Python



Stacking Ensemble Machine Learning With Python



How to Develop Super Learner Ensembles in Python



How to Develop Voting Ensembles With Python



One-vs-Rest and One-vs-One for Multi-Class Classification

How to Develop a Random Forest Ensemble in Python

Photo by Sheila Sund, some rights reserved.

Tutorial Overview

Loving the Tutorials?

This Tutorial is divided into 4 parts, the first part is where you'll find the **Really Good** stuff.

1. Random Forest Algorithm
2. Random Forest for Classification
 - 1. Random Forest for Classification
 - 2. Random Forest for Regression
3. Random Forest Hyperparameters
 - 1. Explore Number of Samples
 - 2. Explore Number of Features
 - 3. Explore Number of Trees
 - 4. Explore Tree Depth
4. Common Questions

Random Forest Algorithm

Random forest is an ensemble of decision tree algorithm.

It is an extension of **bootstrap aggregation (bagging)** and regression problems.

Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

In bagging, a number of decision trees are created where each tree is created from a different bootstrap sample of the training dataset. A **bootstrap sample** is a sample of the training dataset where a sample is drawn with replacement. This is referred to as **sampling with replacement**.



Bagging is an effective ensemble algorithm as each decision tree is fit on a slightly different training sample. Each tree, therefore, has a slightly different performance. Unlike normal decision tree models, such as **classification and regression trees (CART)**, trees used in the ensemble are unpruned, making them more likely to overfit to the training dataset. This is desirable as it helps to make each tree more different and thus less correlated predictions or prediction errors.



[How to Develop Multi-Output Regression Models with Python](#)

Predictions from the trees are averaged across all decision trees resulting in better performance than a single tree in the model.



[Stacking Ensemble Machine Learning With Python](#)



Each model in the ensemble is then used to generate a prediction for a new sample and these predictions are averaged to give the forest's prediction

[How to Develop Super Learner Ensembles in Python](#)

— Page 199, [Applied Predictive Modeling](#), 2013.



[How to Develop Voting Ensembles With Python](#)

Prediction on a regression problem is the average of the prediction across the trees in the ensemble. A prediction on a classification problem is the majority vote for the class label across the trees in the ensemble.



[One-vs-Rest and One-vs-One for Multi-Class Classification](#)

Regression: Prediction is the average prediction across the decision trees.

Classification: Prediction is the majority vote class label predicted across the decision trees.



As with bagging, each tree in the forest casts a vote for the classification of a new sample, and the proportion of votes in each class across the ensemble is the predicted probability

The [Ensemble Learning With Python](#) EBook is where you'll find the **Really Good** stuff.

— Page 3, [Ensemble Learning With Python](#), 2013.

>> SEE WHAT'S INSIDE

Random forest involves constructing a large number of decision trees from bootstrap samples from the training dataset, like bagging.

Unlike bagging, random forest also involves selecting a random subset of features at each split point in the construction of trees. Typically, the value for each input variable in the data in order to select the random subset that may be considered at each split point is to be more different.



Random forests provide an improvement over bagging as they decorrelates the trees. [...] But when building a random forest, a random sample of m predictors is considered, a random sample of m predictors of p predictors.

— Page 320, [An Introduction to Statistical Learning](#), 2013.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

The effect is that the predictions, and in turn, prediction errors, made by each tree in the ensemble are more different or less correlated. When the predictions from these less correlated trees are averaged to make a prediction, it often results in better performance than bagged decision trees.




Perhaps the most important hyperparameter to tune for the random forest is the number of random features to consider at each split point.



[How to Develop Multi-Output Regression Models with Python](#)

Random forests' tuning parameter is the number of randomly selected predictors, k , to choose from at each split, and is commonly referred to as $mtry$. In the regression context, Breiman (2001) recommends setting $mtry$ to be one-third of the number of predictors.



[Stacking Ensemble Machine Learning With Python](#)


— Page 199, [Applied Predictive Modeling](#), 2013.

A good heuristic for regression is to set this hyperparameter to 1/3 the number of input features.



[How to Develop Super Learner Ensembles in Python](#)

$\text{num_features_for_split} = \text{total_input_features} / 3$



For classification problems, Breiman (2001) recommends setting $mtry$ to the square root of the number of predictors.

[How to Develop Voting Ensembles With Python](#)

— Page 387, [Applied Predictive Modeling](#), 2013.



[One-vs-Rest and One-vs-One for Multi-Class Classification](#)

A heuristic for classification is to set this hyperparameter to the square root of the number of input features.

- $\text{num_features_for_split} = \text{sqrt}(\text{total_input_features})$

Loving the Tutorials?

Another important hyperparameter to tune is the depth of the decision trees. Deeper trees are often more overfit to the training data but also more powerful. Deeper, less correlated, which in turn may improve the performance of the ensemble. Depths from 1 to 10 levels may be effective.

>> SEE WHAT'S INSIDE

Finally, the number of decision trees in the ensemble can be set. Often, this is increased until no further improvement is seen.

Want to Get Started With Machine Learning?

Take my free 7-day email crash course.

Click to sign-up and also get a free PDF.

Download Your Free PDF

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Random Forest Scikit-Learn

Random Forest ensembles can be implemented from scratch, although this can be challenging for beginners.



The scikit-learn Python machine learning library provides an implementation of Random Forest for machine learning.

Picked for you:

It is available in modern versions of the library.



How to Develop Multi-Output Regression

Model with Python using a modern version of the library by running the following script:


```
1 # check scikit-learn version
2 import sklearn
3 print(sklearn.__version__)
```

Running the script will print your version of scikit-learn.



Your version should be the same or higher. If not, you must upgrade your version of the scikit-learn library.

```
1 0.22.1
```



How to Develop Voting Ensembles With Python

Random Forest is provided via the `RandomForestRegressor` and `RandomForestClassifier` classes.

Both models operate the same way and take the same arguments that influence how the decision trees are created.



One-vs-Rest and One-vs-One for Multi-Class Classification

Randomness is used in the construction of the model. This means that each time the algorithm is run on the same data, it will produce a slightly different model.

When using machine learning algorithms that have a stochastic learning algorithm, it is good practice to evaluate them by averaging their performance across multiple runs or repeats of **cross-validation**. When fitting a final model, it may be desirable to either increase the number of trees until the variance of the model is reduced across repeated evaluations, or to fit multiple final models and average their prediction.

>> SEE WHAT'S INSIDE

Let's take a look at how to develop a Random Forest ensemble for both classification and regression tasks.

Random Forest for Classification

In this section, we will look at using Random Forest for classification tasks.

First, we can use the `make_classification()` function to generate a synthetic dataset with 1,000 examples and 20 input features.

The complete example is listed below.

```
1 # test classification dataset
2 from sklearn.datasets import make_classification
3 # define dataset
4 X, y = make_classification(n_samples=1000,
5 # summarize the dataset
6 print(X.shape, y.shape)
```

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Running the example creates the dataset and summarizes the shape of the input and output components.



Next, we can evaluate a random forest algorithm on this dataset.

Picked for you:

We will evaluate the model using **repeated stratified k-fold cross-validation**, with three repeats and 10 folds. We will report the mean and standard deviation of the accuracy of the model across all repeats and folds.

```
1 # evaluate random forest algorithm for classification
2 from numpy import mean
3 from numpy import std
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedStratifiedKFold
7 from sklearn.ensemble import RandomForestClassifier
8 # define dataset
9 X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, n_redundant=5,
10 # define the model
11 model = RandomForestClassifier()
12 # evaluate the model
13 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
14 n_scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1, error_score='raise')
15 # report performance
16 print('Accuracy: %.3f (%.3f)' % (mean(n_scores), std(n_scores)))
```



One-vs-Rest and One-vs-One for Multi-Class Classification

Running the example reports the mean and standard deviation accuracy of the model.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

Loving the Tutorials?

In this ensemble, the random forest ensemble with default hyperparameters achieves a classification accuracy of about 90.9 percent.

```
1 Accuracy: 0.909 (0.025)
```

We can also use the random forest model as a final model and make predictions for classification.

First, the random forest ensemble is fit on all available data, then the `predict()` function can be called to make predictions on new data.

The example below demonstrates this on our binary dataset.

```
1 # make predictions using random forest for
2 from sklearn.datasets import make_classification
3 from sklearn.ensemble import RandomForestClassifier
4 # define dataset
5 X, y = make_classification(n_samples=1000,
6 # define the model
7 model = RandomForestClassifier()
8 # fit the model on the whole dataset
9 model.fit(X, y)
10 # make a single prediction
11 row = [-8.52381793, 5.24451077, -12.1496770]
12 yhat = model.predict(row)
13 print('Predicted Class: %d' % yhat[0])
```

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Running the example fits the random forest ensemble model on the entire dataset and is then used to make a prediction on a new row of data, as we might when using the model in an application.



1 Predicted Class: 0

Now that we are familiar with using random forest for classification, let's look at the API for regression.
Picked for you:

Random Forest for Regression




How to Develop Multi-Output Regression

Models with Python

In this section, we will look at using random forests for a regression problem.


First, we can use the `make_regression()` function to create a synthetic regression problem with 1,000 samples and 20 input features.



Stacking Ensemble Machine Learning With Python

The complete example is listed below.

```
1 # test regression dataset
2 from sklearn.datasets import make_regression
3 # define dataset
4 X, y = make_regression(n_samples=1000, n_features=20, n_informative=15, noise=0.1, random_state=42)
5 # summarize the dataset
6 print(X.shape, y.shape)
```



Python

Running the example creates the dataset and summarizes the shape of the input and output components.



One-vs-Rest and One-vs-One for Multi-Class Classification

Next, we can evaluate a random forest algorithm on this dataset.

As we did with the last section, we will evaluate the model using repeated k-fold cross-validation, with three repeats and 10 folds. We will report the mean absolute error (MAE) of the model across all repeats and folds. The `scikit-learn` library makes the MAE negative so that it is maximized instead of minimized. This means that larger negative MAE are better and a perfect model has a MAE of 0.

The complete example is listed below.

```
1 # evaluate random forest ensemble for regression
2 from numpy import mean
3 from numpy import std
4 from sklearn.datasets import make_regression
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedKFold
7 from sklearn.ensemble import RandomForestRegressor
8 # define dataset
9 X, y = make_regression(n_samples=1000, n_features=20, n_informative=15, noise=0.1, random_state=42)
10 # define the model
11 model = RandomForestRegressor()
12 # evaluate the model
13 cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
14 n_scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error', cv=cv)
15 # report performance
16 print('MAE: %.3f (%.3f)' % (mean(n_scores), std(n_scores)))
```

Running the example reports the mean and standard deviation of the MAE across all repeats and folds.

Note: Your results may vary given the stochastic nature of the learning process. Consider running the example multiple times to see the differences in numerical precision. Consider running the example with a different random state to see the differences in numerical precision.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

outcome

Never miss a tutorial:

In this case, we can see the random forest ensemble with default hyperparameters achieves a MAE of about 90.149.

1 MAE: 90.149 (7.924)

We can also use the random forest model as a final model and make predictions for regression.

[How to Develop Multi-Output Regression](#)

[Models with Python](#)

The random forest ensemble is fit on all available data, then the `predict()` function can be called to make predictions on new data.



[Stacking Ensemble Machine Learning](#)

Example below demonstrates this on our regression dataset.

[With Python](#)

```
1 # random forest for making predictions for regression
2 from sklearn.datasets import make_regression
3 from sklearn.ensemble import RandomForestRegressor
4 # define dataset
5 X, y = make_regression(n_samples=1000, n_features=20, n_informative=15, noise=0.1, random_s
6 # define the model
7 model = RandomForestRegressor()
8 # fit the model on the whole dataset
9 model.fit(X, y)
10 # make a single prediction
11 row = [[-0.89483109, -1.0670149, -0.25448694, -0.53850126, 0.21082105, 1.37435592, 0.71203659, 0.7
12 yhat = model.predict(row)
13 print('Prediction: %d' % yhat[0])
```



[One-vs-Rest and One-vs-One for Multi-](#)

[Class Classification](#)

The example fits the random forest ensemble model on the entire dataset and is then used to make a prediction on a new row of data, as we might when using the model in an application.

1 Prediction: -173

Loving the Tutorials?

Now that we are familiar with using the scikit-learn API to evaluate and use random forest ensembles,

let's [The Fast and Easy Learning With Python EBook](#)

is where you'll find the **Really Good** stuff.

Random Forest Hyperparameters

>> SEE WHAT'S INSIDE

In this section, we will take a closer look at some of the hyperparameters you should consider tuning for the random forest ensemble and their effect on model performance.

Explore Number of Samples

Each decision tree in the ensemble is fit on a bootstrap sample of the training dataset.

This can be turned off by setting the “*bootstrap*” argument to `False`. If `True`, a bootstrap sample of the training dataset will be used to train each decision tree.

The “*max_samples*” argument can be set to a float between 0 and 1, representing the fraction of the training dataset to make the bootstrap sample.

For example, if the training dataset has 100 rows, and `max_samples` is set to 0.5, then each decision tree will be fit on a bootstrap sample of 50 rows.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

A smaller sample size will make trees more different, and a larger sample size will make the trees more similar. Setting `max_samples` to "None" will make the sample size the same size as the training dataset



The example below demonstrates the effect of different bootstrap sample sizes from 10 percent to 100 percent on the random forest algorithm.

```

1 # explore to develop further about bootstrap sample size on performance
2 from numpy import mean
3 from numpy import std
4 from numpy import arange
5 from sklearn.datasets import make_classification
6 from sklearn.model_selection import cross_val_score
7 from sklearn.model_selection import RepeatedStratifiedKFold
8 from sklearn.ensemble import RandomForestClassifier
9 from matplotlib import pyplot
10
11 # get the dataset
12 def get_dataset():
13     # make a Python classification
14     return X, y
15
16 # get a list of models to evaluate
17 def get_models():
18     models = dict()
19     # explore ratios from 10% to 100% in 10% increments
20     for i in arange(0.1, 1.1, 0.1):
21         key = '%.1f' % i
22         # set max_samples=None to use 100%
23         if i == 1.0:
24             i = None
25         models[key] = RandomForestClassifier(max_samples=i)
26     return models
27
28 # evaluate a given model using cross-validation
29 def evaluate_model(model, X, y):
30     # define the evaluation procedure
31     # The Ensemble Repeatedly With PythoRFBK
32     # evaluate the model and collect the results
33     scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
34     return scores
35     >> SEE WHAT'S INSIDE
36
37 # define dataset
38 X, y = get_dataset()
39 # get the models to evaluate
40 models = get_models()
41 # evaluate the models and store results
42 results, names = list(), list()
43 for name, model in models.items():
44     # evaluate the model
45     scores = evaluate_model(model, X, y)
46     # store the results
47     results.append(scores)
48     names.append(name)
49     # summarize the performance along the
50     print('>%s %.3f (%.3f)' % (name, mean(scores), std(scores)))
51 # plot model performance for comparison
52 pyplot.boxplot(results, labels=names, show=True)
53 pyplot.show()

```

Running the example first reports the mean accuracy for each model.

Note: Your results may vary given the stochastic nature of the random forest algorithm. Consider running the example multiple times to see the differences in numerical precision. Consider running the example multiple times to see the outcome.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

In this case, the results suggest that using a bootstrap sample size that is equal to the size of the training dataset achieves the best results on this dataset.



This is the default and should probably be used in most cases.

```
1 >10 0.856 (0.031)
2 >20 0.873 (0.029)
3 >30 0.881 (0.021)
4 >40 0.891 (0.033)
5 >50 0.893 (0.025)
6 >60 0.897 (0.030)
7 >70 0.902 (0.024)
8 >80 0.903 (0.024)
9 >90 0.900 (0.026)
10 >100 0.903 (0.027)
```


A box and whisker plot is created for the distribution of accuracy scores for each bootstrap sample size.



How to Develop Supercharged
Ensembles in Python

In this case, we can see a general trend that the larger the sample, the better the performance of the

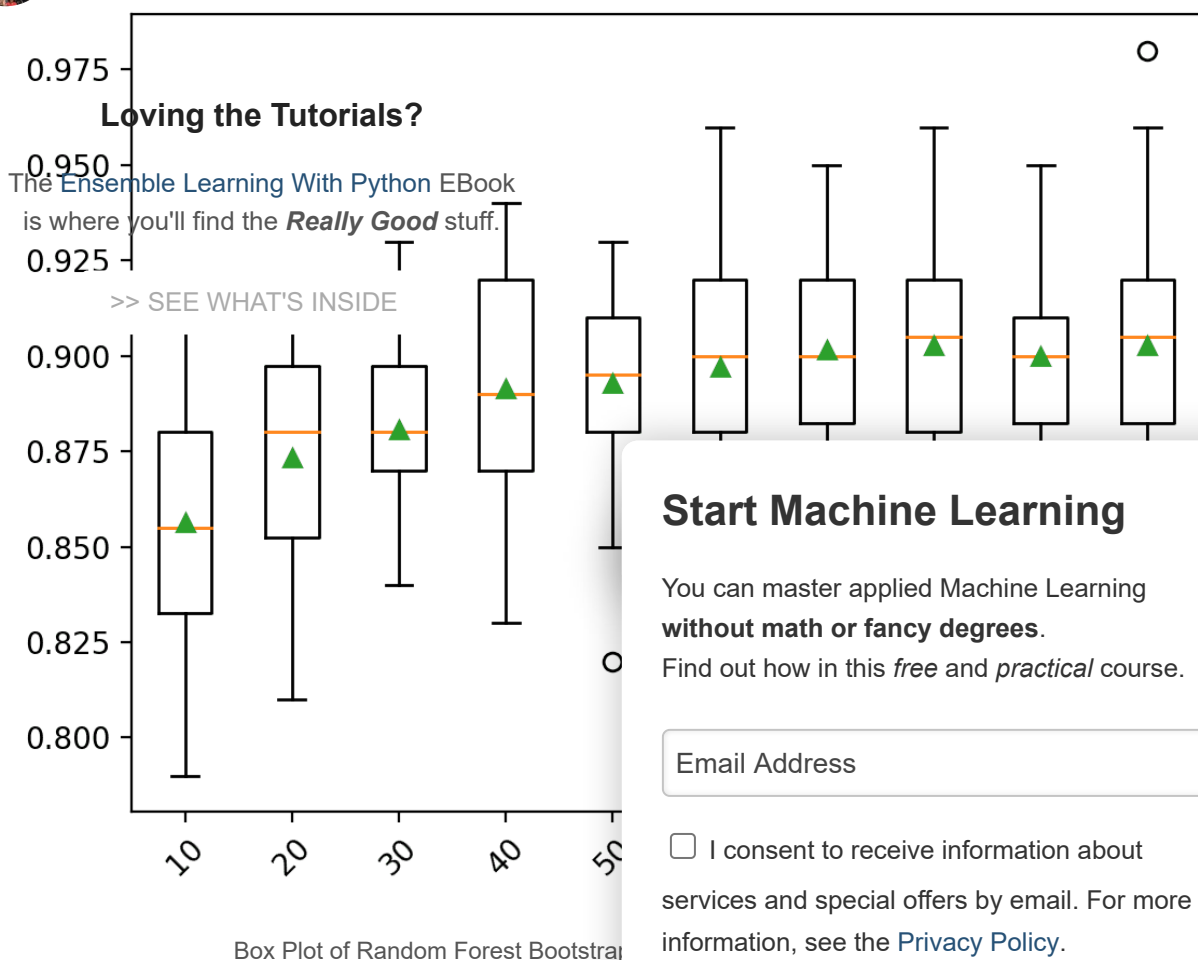
You might like to extend this example and see what happens if the bootstrap sample size is larger or much larger than the training dataset (e.g. you can set an integer value as the number of samples of a float percentage of the training dataset size).



How to Develop Voting Ensembles With
Python



One-vs-Rest and One-vs-One for Multi-
Class Classification



Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Explore Number of Features

The number of features that is randomly sampled for each split point is perhaps the most important feature to configure for random forest.

It is set via the `max_features` argument and defaults to the square root of the number of input features. In this case, for our test dataset, this would be $\sqrt{20}$ or about four features.



How to Develop Multi-Output Regression

Model below explores the effect of the number of features randomly selected at each split point on model accuracy. We will try values from 1 to 7 and would expect a small value, around four, to perform well based on the heuristic.

```
Stacking Ensemble Machine Learning
With Python
1 # explore random forest number of features effect on performance
2 from numpy import mean
3 from numpy import std
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedStratifiedKFold
7 from sklearn.ensemble import RandomForestClassifier
8 from matplotlib import pyplot
9
10 # get the dataset
11 def get_dataset():
12     X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, n_redundant=5)
13     return X, y
14
15 # get a list of models to evaluate
16 def get_models():
17     models = []
18     # explore number of features from 1 to 7
19     for i in range(1, 8):
20         models.append(RandomForestClassifier(max_features=i))
21     return models
22
23 # evaluate a given model using cross-validation
24 def evaluate_model(model, X, y):
25     # define the evaluation procedure
26     cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
27     # evaluate the model and collect the results
28     scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
29     return scores
30
31 # define dataset
32 X, y = get_dataset()
33 # get the models to evaluate
34 models = get_models()
35 # evaluate the models and store results
36 results, names = list(), list()
37 for name, model in models.items():
38     # evaluate the model
39     scores = evaluate_model(model, X, y)
40     # store the results
41     results.append(scores)
42     names.append(name)
43     # summarize the performance along the
44     print('>%s %.3f (%.3f)' % (name, mean(scores), std(scores)))
45 # plot model performance for comparison
46 pyplot.boxplot(results, labels=names, show=False)
47 pyplot.show()
```

Running the example first reports the mean accuracy

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average



In this case, the results suggest that a value between three and five would be appropriate, confirming **Picked for you** result of four on this dataset. A value of five might even be better given the smaller standard deviation in classification accuracy as compared to a value of three or four.



[How to Develop Multi-Output Regression](#)

```
1 >1 0.897 (0.025)
2 >2 0.900 (0.028)
3 >3 0.903 (0.027)
4 >4 0.903 (0.022)
5 >5 0.903 (0.019)
6 >6 0.898 (0.025)
7 >7 0.900 (0.024)
```


A box and whisker plot is created for the distribution of accuracy scores for each feature set size.



[How to Develop Super Learner](#)

[Ensembles in Python](#)

I see a trend in performance rising and peaking with values between three and five and falling again as larger feature set sizes are considered.

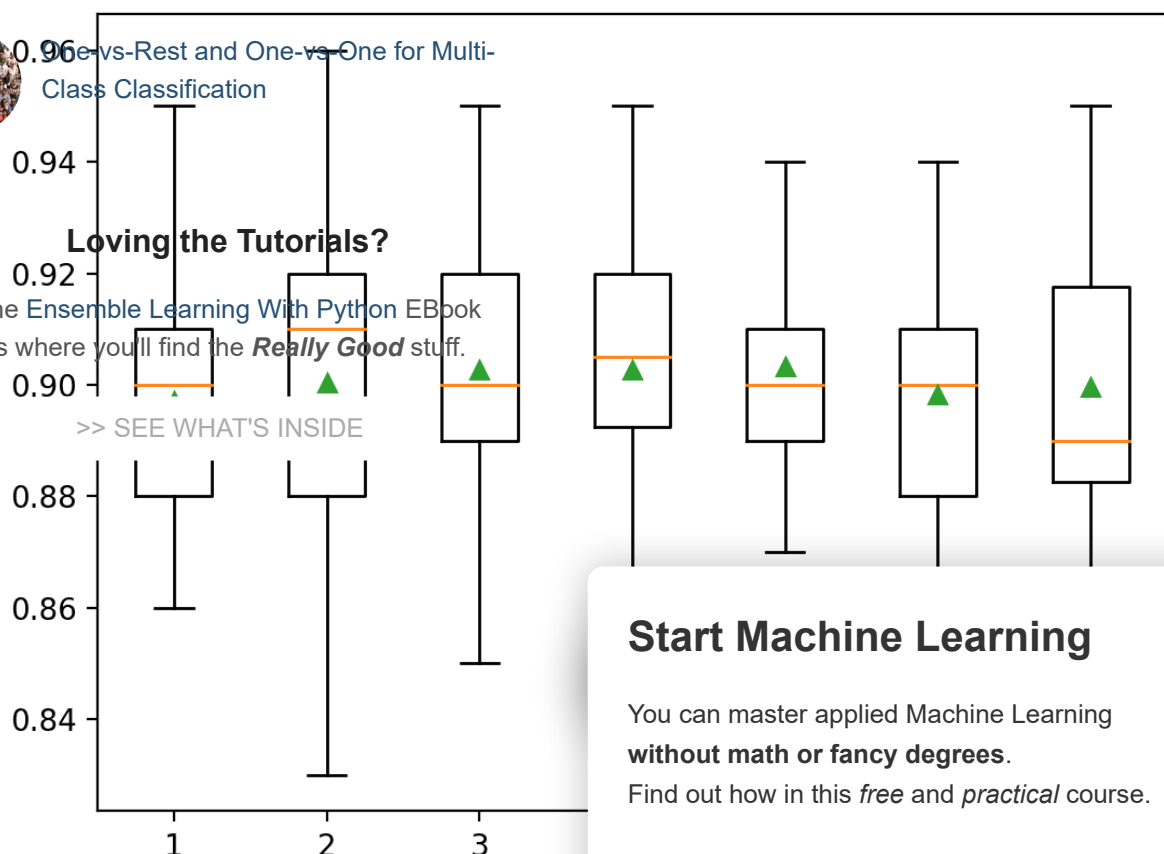


[How to Develop Voting Ensembles With Python](#)



[One-vs-Rest and One-vs-One for Multi-Class Classification](#)

Loving the Tutorials?
The [Ensemble Learning With Python EBook](#) is where you'll find the *Really Good* stuff.
>> SEE WHAT'S INSIDE



Box Plot of Random Forest Feature

Explore Number of Trees

The number of trees is another key hyperparameter

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free and practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

Typically, the number of trees is increased until the model performance stabilizes. Intuition might suggest that more trees will lead to overfitting, although this is not the case. Both bagging and random forest algorithms are somewhat immune to overfitting the training dataset given the stochastic nature of the learning algorithm.

Picked for you: The number of trees can be set via the “*n_estimators*” argument and defaults to 100.

How to Develop Multi-Output Regression Models with Python This sample code explores the effect of the number of trees with values between 10 to 1,000.

```
1 # explore random forest number of trees effect on performance
2 from numpy import mean
3 from numpy import std
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedStratifiedKFold
7 from sklearn.ensemble import RandomForestClassifier
8 from matplotlib import pyplot
9
10 # get the dataset
11 def get_dataset():
12     X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, n_redundant=5)
13     return X, y
14
15 # get a list of models to evaluate
16 def get_models():
17     models = dict()
18     # define number of trees to consider
19     n_trees = [10, 50, 100, 500, 1000]
20     for n in n_trees:
21         models[n] = RandomForestClassifier(n_estimators=n)
22     return models
23
24 # evaluate a given model using cross-validation
25 def evaluate_model(model, X, y):
26     # define the evaluation procedure
27     cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
28     # evaluate the model and collect the results
29     scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
30     return scores
31
32 # define dataset
33 X, y = get_dataset()
34 # get the models to evaluate
35 models = get_models()
36 # evaluate the models and store results
37 results, names = list(), list()
38 for name, model in models.items():
39     # evaluate the model
40     scores = evaluate_model(model, X, y)
41     # store the results
42     results.append(scores)
43     names.append(name)
44     # summarize the performance along the way
45     print('>%s %.3f (%.3f)' % (name, mean(scores), std(scores)))
46 # plot model performance for comparison
47 pyplot.boxplot(results, labels=names, show=False)
48 pyplot.show()
```

Running the example first reports the mean accuracy for each number of trees.

Note: Your results may vary given the stochastic nature of the learning algorithm. Differences in numerical precision. Consider running the example multiple times to see the outcome.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

In this case, we can see that performance rises and stays flat after about 100 trees. Mean accuracy scores fluctuate across 100, 500, and 1,000 trees and this may be statistical noise.

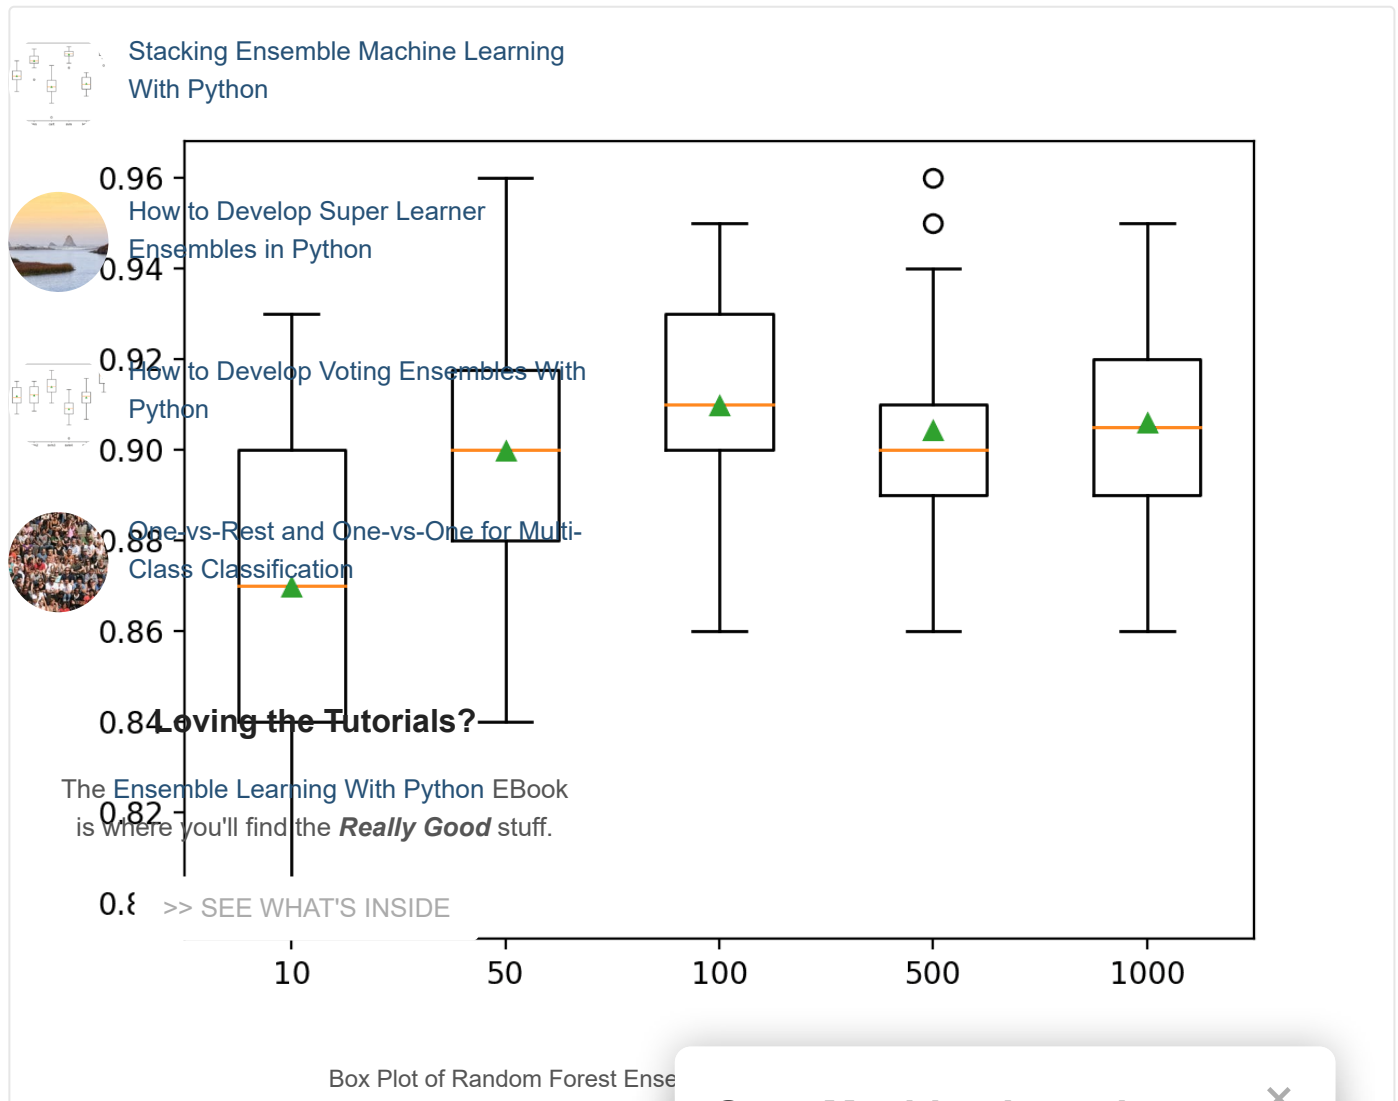


- 1 >10 0.870 (0.036)
- 2 >50 0.900 (0.028)
- 3 >100 0.910 (0.024)
- 4 >500 0.904 (0.024)
- 5 >1000 0.906 (0.023)



How to Develop Multi-Output Regression Models with Python

and whisker plot is created for the distribution of accuracy scores for each configured number of



Explore Tree Depth

A final interesting hyperparameter is the maximum

By default, trees are constructed to an arbitrary depth although we can also explore fitting trees with different

The maximum tree depth can be specified via the `max_depth` parameter (the default is `None`, which means no limit on the maximum depth) by default.

The example below explores the effect of random

1 # explore random forest tree depth effect

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

```

2 from numpy import mean
3 from numpy import std
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedStratifiedKFold
7 from sklearn.ensemble import RandomForestClassifier
8 from matplotlib import pyplot
9
10 # get the dataset
11 def get_dataset():
12     X, y = make_classification(n_samples=1000, n_features=20, n_informative=15, n_redundant=5,
13                               random_state=1)
14     return X, y
15
16 # get a list of models to evaluate
17 def get_models():
18     models = dict()
19     # consider tree depths from 1 to 7 and None=full
20     depths = [i for i in range(1,8)] + [None]
21     for n in depths:
22         models[str(n)] = RandomForestClassifier(max_depth=n)
23     return models
24
25 # evaluate a given model using cross-validation
26 def evaluate_model(model, X, y):
27     # define the evaluation procedure
28     cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
29     # evaluate the model and collect the results
30     scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
31     return scores
32
33 # define dataset
34 X, y = get_dataset()
35 # get the models to evaluate
36 models = get_models()
37 # evaluate the models and store results
38 results, names = list(), list()
39 for name, model in models.items():
40     # evaluate the model
41     scores = evaluate_model(model, X, y)
42     # store the results
43     results.append(scores)
44     names.append(name)
45 # summarize the performance along the way
46 print('>%s %.3f (%.3f)' % (name, mean(scores), std(scores)))
47 # plot model performance for comparison
48 pyplot.boxplot(results, labels=names, showmeans=True)
49 pyplot.show()

```

Running the example first reports the mean accuracy for each configured maximum tree depth.

Note: Your results may vary given the stochastic nature of the data and differences in numerical precision. Consider running the example multiple times to get a better outcome.

In this case, we can see that larger depth results in better performance, with the maximum depth achieving the best performance of 0.903.

```

1 >1 0.771 (0.040)
2 >2 0.807 (0.037)
3 >3 0.834 (0.034)
4 >4 0.857 (0.030)
5 >5 0.872 (0.025)
6 >6 0.887 (0.024)
7 >7 0.890 (0.025)
8 >None 0.903 (0.027)

```

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

A box and whisker plot is created for the distribution of accuracy scores for each configured maximum tree depth.




In this case, we can see a trend of improved performance with increase in tree depth, supporting the default of no maximum depth.

Picked for you:




How to Develop Multi-Output Regression Models with Python



Stacking Ensemble Machine Learning With Python



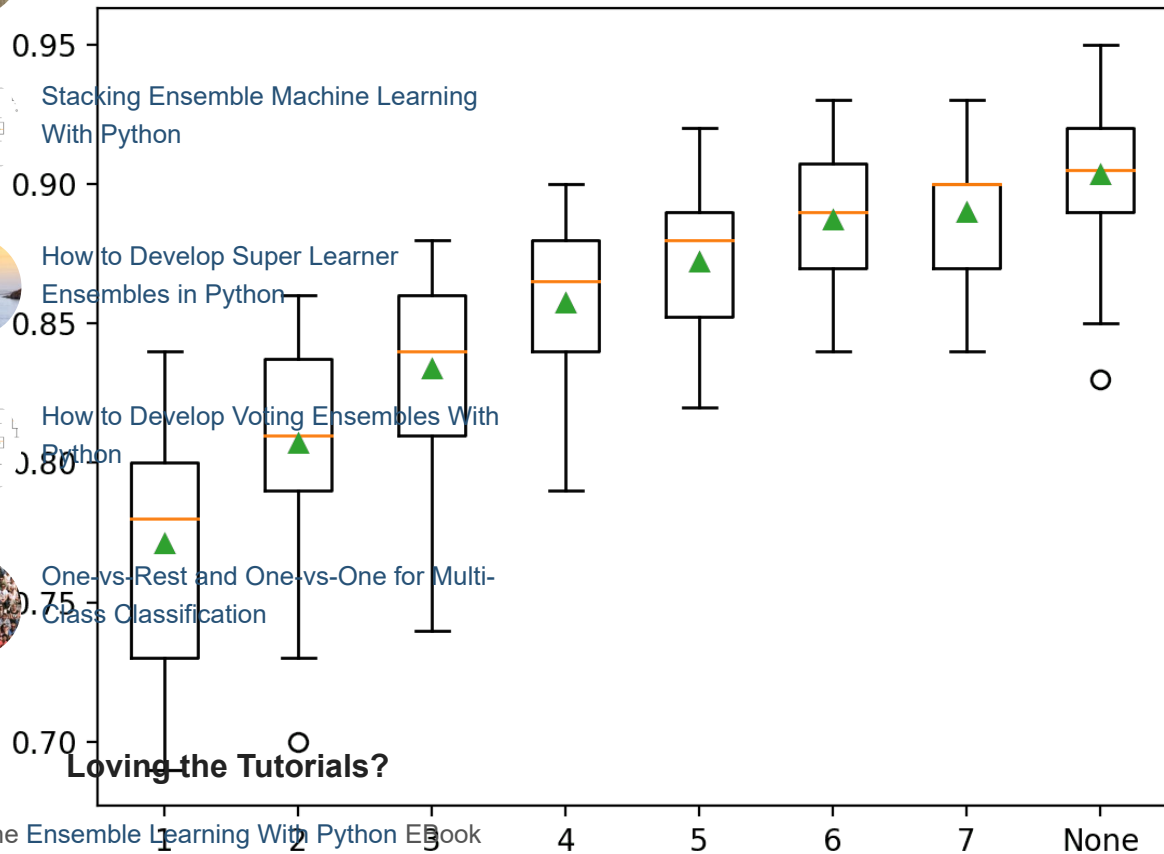
How to Develop Super Learner Ensembles in Python



How to Develop Voting Ensembles With Python



One-vs-Rest and One-vs-One for Multi-Class Classification



Loving the Tutorials?

The Ensemble Learning With Python EBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE >> Forest Maximum Tree Depth vs. Classification Accuracy

Common Questions

In this section we will take a closer look at some of the common questions about the random forest ensemble procedure.

Q. What algorithm should be used in the ensemble?

Random forest is designed to be an ensemble of decision trees.

Q. How many ensemble members should be used?

The number of trees should be increased until no further improvement is seen on the dataset.



As a starting point, we suggest using at least 100 trees. In many cases, profiles are still improving at 1,000 trees, though.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial:



Q. Won't the ensemble overfit with too many trees? Picked for you:


No. Random forest ensembles (do not) are very unlikely to overfit in general.



[How to Develop Multi-Output Regression](#)

[Models with Python](#)

Another claim is that random forests “cannot overfit” the data. It is certainly true that increasing [the number of trees] does not cause the random forest sequence to overfit ...



[Stacking Ensemble Machine Learning](#)

[With Python](#)

Page 596, [The Elements of Statistical Learning](#), 2016.

Q. How large should the bootstrap sample be?




[How to Develop Super Learner](#)

[Ensembles in Python](#)

Good practice is to make the bootstrap sample as large as the original dataset size.

That is 100% the size or an equal number of rows as the original dataset.



[How to Develop Voting Ensembles With](#)

[Python](#)

How many features should be chosen at each split point?

The best practice is to test a suite of different values and discover what works best for your dataset.



[One-vs-Rest and One-vs-One for Multi-](#)

[Class Classification](#)

As a heuristic, you can use:

- **Classification:** Square root of the number of features.
- **Regression:** One third of the number of features.

Loving the Tutorials?

Q. What problems are well suited to random forest?

is where you'll find the **Really Good** stuff.

Random forest is known to work well or even best on a wide range of classification and regression problems. >> SEE WHAT'S INSIDE



The authors make grand claims about the success of random forests: “most accurate”, “most interpretable”, and the like. In our experience, random forests are “most accurate” with little tuning required.

— Page 590, [The Elements of Statistical Learning](#)

Further Reading

This section provides more resources on the topic

Tutorials

- [How to Implement Random Forest From Scratch](#)

Papers

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

• [Random Forests](#), 2001.

Never miss a tutorial:

Books



• [Applied Predictive Modeling](#), 2013.

• [The Elements of Statistical Learning](#), 2016.


Picked for you:

• [An Introduction to Statistical Learning with Applications in R](#), 2014.



[How to Develop Multi-Output Regression Models with Python](#)

• [sklearn.ensemble.RandomForestRegressor API](#).



[sklearn.ensemble.RandomForestClassifier API](#).

[Stacking Ensemble Machine Learning With Python](#)


Articles



[Random Forest, Wikipedia](#)

[How to Develop Super Learner Ensembles in Python](#)

Summary



In this tutorial, you discovered how to develop random forest ensembles for classification and regression with Python.

Specifically, you learned:



• [One-vs-Rest and One-vs-One for Multi-Class Classification](#)

• A random forest ensemble is an ensemble of decision trees and a natural extension of bagging.

- How to use the random forest ensemble for classification and regression with scikit-learn.
- How to explore the effect of random forest model hyperparameters on model performance.

Loving the Tutorials?


Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

The [Ensemble Learning With Python EBook](#) is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE

Get a Handle on Modern Ensemble Learning!



Improve Your Predictions in Minutes

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial:




More On This Topic


Picked for you:



How to Develop Multi-Output Regression Models with Python




Stacking Ensemble Machine Learning With Python




Bagging and Random Forest Ensemble Algorithms for...



How to Develop Super Learner Ensembles in Python



How to Develop Voting Ensembles With Python




One-vs-Rest and One-vs-One for Multi-Class Classification

How to Develop Random Forest Ensembles With XGBoost

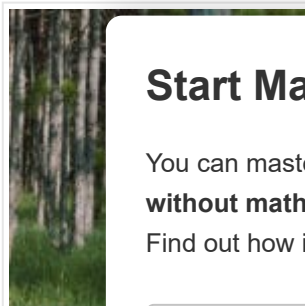
Loving the Tutorials?

The [Ensemble Learning With Python](#) EBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



How to Develop a Random Subspace Ensemble With Python



How to Implement Random Subspace Ensemble With Python

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).


START MY EMAIL COURSE

Never miss a tutorial:**Picked for you:**

How to Develop Multi-Output Regression Models with Python




Use Random Forest: Testing 179 Classifiers on 121 Datasets




Stacking Ensemble Machine Learning With Python



How to Develop Super Learner Ensembles in Python



Tune Machine Learning Algorithms in R (random forest...)



How to Develop Voting Ensembles With Python

**About Jason Brownlee**

One-vs-Rest and One-vs-One for Multi-Class Classification

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

Loving the Tutorials?

< [How to Develop Voting Ensembles With Python](#)
The Ensemble Learning With Python EBook is where you'll find the **Really Good** stuff.

[How to Develop an Extra Trees Ensemble with Python](#) >

>> SEE WHAT'S INSIDE

36 Responses to *How to Develop a Random Forest Ensemble in Python*



dcart April 20, 2020 at 6:20 am #

Hi Jason

Hope you are doing well in this time of lock down.

Great articles as usual. Although I like to apply RF memory as the data is huge.

Any advise using RF for huge data set?

Thanks

Dennis

Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial:**Jason Brownlee**

April 20, 2020 at 7:36 am #


REPLY ↩



prototype on a small sample of data first to see if it is effective.

More ideas here:

Picked for you:
<https://machinelearningmastery.com/faq/single-faq/how-can-i-run-large-models-or-models-on-lots-of-data>

How to Develop Multi-Output Regression Models with Python

Stacking Ensemble Machine Learning With Python


Hi Jason, Are you planning a new book on EnsembleS?

REPLY ↩


How to Develop Super Learner Ensembles in Python

April 23, 2020 at 6:12 am #

REPLY ↩


How to Develop Voting Ensembles With Python

Not sure at this stage.

Are there ensemble topics you'd like me to write about?


One-vs-Rest and One-vs-One for Multi-Class Classification

April 24, 2020 at 5:17 am #

REPLY ↩

probably advanced stacking and how to win kaggle/data science competitions

Loving the Tutorials?The **Ensemble Learning With Python EBook**is where you'll find the **Really Good** stuff.

April 24, 2020 at 5:54 am #

REPLY ↩

<https://machinelearningmastery.com/stacking-ensemble-machine-learning-with-python/>

And this:

<https://machinelearningmastery.com/super-learner-ensemble-in-python/>
**Stelios** April 26, 2020 at 4:34 pm #

Thank you so much for your great support

**Jason Brownlee** April 27, 2020 at 5:30 am #

You're welcome.

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial:**David Sanchez** May 10, 2020 at 11:25 pm #

REPLY ↩




Hi Jason,

I'm implementing a Random Forest and I'm getting a shifted time-series in the predictions. If I build the model for predicting e.g. 4 steps ahead, my time-series of predictions seems 4 steps shifted to the right comparing to my time-series of observations. If I try to predict 16 steps ahead, it seems 16 steps shifted. How to Develop Multi-Output Regression doesn't seem to be happening?



Thanks for all your tutorials!



[Stacking Ensemble Machine Learning With Python](#)


Jason Brownlee May 11, 2020 at 6:00 am #

REPLY ↩



[How to Develop Super Learner Ensembles in Python](#)

Yes, it sounds like the model has learned a persistence (no skill) forecast. E.g. it predicts the input as the output.



[How to Develop Voting Ensembles With Python](#)

David Sanchez May 14, 2020 at 6:28 pm #

REPLY ↩



Hi Jason,
[One-vs-Rest and One-vs-One for Multi-Class Classification](#)

Thank you for your reply.

I'm also wondering, if I try to build a model where my train set has more variables than my test set, how should I proceed?

Loving the Tutorials?

As far as I've seen about it, I should recreate those missing variables in my test dataframe and set them as 0.

The [Ensemble Learning With Python](#) EBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE

**Jason Brownlee** May 15, 2020 at 5:57 am #

REPLY ↩

The number of variables (columns) must be the same in train and test sets.

**Grzegorz Kępiś** May 26, 2020 at 10:46 pm

Very nice tutorial of RF usage!

It is really practical to know good practices on those very competitive in real industrial applications! (often Networks).

Regards!

**Jason Brownlee** May 27, 2020 at 7:54 am**Start Machine Learning** ✕

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Thanks!
Never miss a tutorial:

Agreed. XGBoost more so.



Picked for you:

May 28, 2020 at 10:03 am #

REPLY ↩



How to Develop Multi-Output Regression

Models with Python

Hi Jason,


use, check it.

"This means that larger negative MAE are better and a perfect model has a MAE of 0."

Thank you! enjoying a lot this stuff.

Stacking Ensemble Machine Learning

With Python



Jason Brownlee

How to Develop Super Learner


Ensembles in Python

It is correct.

-10 is greater than -100.

How to Develop Voting Ensembles With

Python



One-vs-Rest and One-vs-One for Multi-

Class Classification

October 16, 2020 at 4:26 pm #

REPLY ↩

Hello Jason, Please I have a question

I have the following situation that is already programmed with Logistic regression, I have tried the same program with Random Forest in order to check how it could improve the accuracy.

Actually, the accuracy was improved, but I don't know if it is logical to use the Random Forest in my

The Ensemble Learning With Python Ebook

is where you'll find the **Really Good** stuff.

My case study is as follow :

Based >> SEE WHAT'S INSIDE to predict if a customer will buy a product or not depending on his prior history. I.e to know now much a customer bought the same product previously, and how much he just check it without buying it

The used data has the following structure:

Id clients CurrentProd P1+ P1- P2+ P2- P3+ P3- ..

10 CL1 P1, P3 6 1 0 0 8 2 0 0 1

11 CL1 P1, P2 7 1 5 2 0 0 0 0 1

with:

CurrentProd: means a list of products that I need to

P1+: mean how many time a client buy product 1,

P1-: refers to the number that a client checked a pr

columns present all products existing in the market

PRODUCT) and at each row the most of those row

CurrentPRod

So I want to know if the random forest could be use

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

PS: I must use the data as it is without any change in features or structure
Never miss a tutorial:



manuela October 16, 2020 at 4:43 pm #

REPLY ↩

Picked for you:

Id..|.clients..|.CurrentProd..|.P1+..|.P1-..|.P2+..|.P2-..|.P3+..|.P3-..|. ...|.PN+..|.PN-..|.Output



10 | CL 1 | P1, P3 | 6 | 1 | 0 | 0 | 8 | 2 | 0 | 0 | 1
 1 | Models with Python | P1, P2 | 7 | 1 | 5 | 2 | 0 | 0 | 0 | 0 | 1



Stacking Ensemble Machine Learning



Jason Brownlee October 17, 2020 at 5:58 am #

REPLY ↩

Perhaps try it and compare results.



How to Develop Super Learner

The key will to find an appropriate representation for the problem. This may give you ideas (replace site with product).

<https://machinelearningmastery.com/faq/single-faq/how-to-develop-forecast-models-for-multiple-sites>



How to Develop Voting Ensembles With Python



manuela October 17, 2020 at 6:45 am #

REPLY ↩



One-vs-Rest and One-vs-One for Multi-Class Classification

Thanks for your quick replay,
 I have already tried it, and it gives me a good result,

but I want to know if it is logical to use it with 200 features (Product1, Product2....)

Loving the Tutorials?

The Ensemble Learning With Python EBook



Jason Brownlee

October 17, 2020 at 1:41 pm #

REPLY ↩

is where you find the **Really Good** stuff.

I use the features that result in the best performance, regardless of how many.

>> SEE WHAT'S INSIDE



dinesh December 21, 2020 at 7:29 pm #

REPLY ↩

how to decide these paramters

n_samples=1000, n_features=20, n_informative=1

any suggestion on this please.



Jason Brownlee December 22, 2020 at 6:00 am #

This defines the test problem, it is con

The ideas is you replace this with your own dat

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.
 Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about
 services and special offers by email. For more
 information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial:**dinesh** December 25, 2020 at 1:51 pm #

REPLY ↩




what is the best practice , also if i want to learn about the meaning of these parameter.
please share some information about the hyper parameter tuning.

Picked for you:**How to Develop Multi-Output Regression****Models With Python** **Jason Brownlee** December 26, 2020 at 5:08 am #

REPLY ↩

Best practice for test problems? I don't understand.

**Stacking Ensemble Machine Learning**
With Python


Best practice for hyperparameter tuning is in the above example, e.g. grid search.

**How to Develop Super Learner**
Ensembles in Python

February 17, 2021 at 6:03 pm #

REPLY ↩

Perfect!

**How to Develop Voting Ensembles With**
Python**Jason Brownlee** February 18, 2021 at 5:12 am #

REPLY ↩

**One-vs-Rest and One-vs-One for Multi-**
Class Classification

Thanks!

Loving the Tutorials?**Alex** September 8, 2021 at 12:29 am #

REPLY ↩

The Ensemble Learning With Python EBook

is where you'll find the **Really Good** stuff.

I wonder to know is there any way to find out that under which condition my model has wrong

>> SEE WHAT'S INSIDE
any way to find (range) values for features that tell me that the
odel is not reliable. Or let machine learn that when the prediction
could not be reliable.

I think it could be, some how , the other way around of machine learning ,isnt it?

Any suggestions for it?

**Adrian Tam** September 8, 2021 at 12:29 am #Quite impossible to know bec
data provided. You're simply asking wh**Start Machine Learning**

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about
services and special offers by email. For more
information, see the [Privacy Policy](#).

START MY EMAIL COURSE

**Francisco Pérez Liébana** March 12, 2021 at 12:29 am #

Hello Jason,

Thanks for your articles, they are very useful!
Never miss a tutorial:

Do you know how can I get a graphic representation of the trees in the trained model ? I was trying to use `export_gaphviz` but using "cross_val_scores" function fitting estimator on its own, i don't know how to use `export_gaphviz` function.

Picked for you:
 Thanks in advance for your answer.
 Francisco



[How to Develop Multi-Output Regression Models with Python](#)



Jason Brownlee March 13, 2021 at 5:33 am #

REPLY ↩

[Checking Ensemble Machine Learning](#)

With Python. I believe it's possible but I have not done it before, sorry Francisco.



[How to Develop Super Learner](#)

Ian Sharler May 12, 2021 at 11:59 pm #

REPLY ↩

Hi Jason,



[How to Develop Voting Ensembles With Python](#)
 Thanks for the clear and useful introduction.

I have a question on how the Random Forest algorithm handles missing features.



For example suppose the data set is a 24H time series, for which I want to build a classifier.

[One-vs-Rest and One-vs-One for Multi-Class Classification](#)

Some of the features are available only in daytime, some only in night-time, and some others are partly unavailable.

What is the best way to adapt the algorithm to address this task.

Loving the Tutorials?

Thanks in advance,

[The Ensemble Learning With Python EBook](#)

Ilan

is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE



Jason Brownlee May 13, 2021 at 6:03 am #

REPLY ↩

Perhaps try a suite of approaches for handling the missing data and discover what works well or best for your dataset.



Ahmad Afif Aulia Hariz July 13, 2021 at 10

Thanks a lot for the article.

I've just build my own RF Regressor, i have (2437, about 0.7

I want to improve it into 0.95. Any suggestion?

Thanks for help!


Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Never miss a tutorial!








Jason Brownlee

July 14, 2021 at 5:29 am #


REPLY

Perhaps some of these ideas will help:



<https://machinelearningmastery.com/machine-learning-performance-improvement-cheat-sheet/>

Picked for you:



ANANTHA KRISHNAN

How to Develop Voting Ensembles With Python Regression Models with Python

July 14, 2021 at 7:41 pm #

REPLY

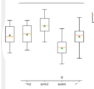
Hai sir,

I m having two seperate data frames. One for Training and another one for Testing. I need to perform Stacking Ensemble Machine Learning With Python Random Forest Classification. How to load these files to Random Forest without splitting.




How to Develop Super Learner Ensembles in Python

See a Reply



How to Develop Voting Ensembles With Python



One-vs-Rest and One-vs-One for Multi-Class Classification

Loving the Tutorials? (required)


The Ensemble Learning With Python EBook is where you'll find the Really Good stuff.

Email (will not be published) (required)

>> SEE WHAT'S INSIDE

Website

SUBMIT COMMENT



Welcome!
I'm Jason Brownlee PhD and I help developers get results with
[Read more](#)

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.** Find out how in this *free* and *practical* course.

Email Address

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE


Never miss a tutorial:



Picked for you:




How to Develop Multi-Output Regression Models with Python



Stacking Ensemble Machine Learning With Python



How to Develop Super Learner Ensembles in Python



How to Develop Voting Ensembles With Python



One-vs-Rest and One-vs-One for Multi-Class Classification

Loving the Tutorials?

The Ensemble Learning With Python EBook is where you'll find the **Really Good** stuff.

>> SEE WHAT'S INSIDE

© 2021 Machine Learning Mastery. All Rights Reserved.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

Start Machine Learning ×

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

☐ I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE